

## CSE 232, Lab Exercise 7

### *Command Line*

### *Process Management (&, bg/fg)*

So far, we've been teaching you how to run programs at the command line. Running the program and waiting for it to finish works okay for short/fast programs, but if you have a program that takes minutes or hours to finish, you don't want to have to wait that long to do other things at the terminal. Instead, there is a way to run the program in the *background*, so that it is running but not blocking your access to the terminal.

### **Executing a Background Job**

Let us say we have a long-running program called `a.out`. Normally you would run the program like so:

```
./a.out
```

But if you did that your terminal would be blocked until it finished. No commands can be executed until it is done. If instead you add an ampersand (&) *after* the command, the execution is run in the background.

```
./a.out &
```

The job, `a.out`, will now run but the terminal is freed for your commands.

### **What jobs do I have?**

You can run the command `jobs` to see the status of your suspended (and running) jobs. It reports it in the following way (I have a job called `top` running in background)

```
>jobs
[1]+  Stopped                  top
```

The `[1]` is your local job number. We will use that in a minute.

### **Suspending a Running Job and Running in the Background**

Lets say you already were running `a.out` in the foreground (without the ampersand). You already know about CTRL-C to kill a running program. Here's a new one, CTRL-Z. CTRL-Z doesn't kill a program, but instead suspends it (pauses it), giving you back control of your terminal. Whatever was running is now stopped, but can be restarted.

After suspending a job, you can have the suspended job run in background using the command `bg`

## Take a Background/Suspended Job and Run It In The Foreground

If you want to see the output from a suspended/background job, you can do so with the `fg` command. By default, `fg` makes the most recent background job run in the foreground. But if you supply the job ID number (from the `jobs` command), you can select a specific background job to be moved to the foreground like so:

```
fg %1
```

## Killing Background Jobs

To kill a specific background job, use the `kill` command and the job id. Like so:

```
kill %2
```

For example, the `top` command tells you about the present status of your computer. It continuously prints info about processes, memory usage etc. I can do the following:



Make a program that outputs all the numbers between 1 and 1 million. Run the program in the background and demonstrate moving it to the foreground and killing it to the TA.

## *The Problem*

We are going to work on 2D vectors.

### **2D vector as a matrix**

You remember matrices, don't you? We are going to do some simple manipulation of a matrix, namely: adding two matrices and multiplying a matrix by a scalar. You watched the 2D vector video, right? RIGHT???

## **Matrix**

A **matrix** is a 2-dimensional (rows and columns) data structure. It has a shape indicated by the number of rows and the number of columns. Though I suppose a matrix could have uneven sized rows, this doesn't usually happen in practice so a matrix is always rectangular, potentially square (based on its shape).

rows				
0	1	2	3	
1	4	5	6	
2	7	8	9	
	0	1	2	
				cols


shape = rows × cols (here 3 × 3)

## Matrix operations


We will perform two operations on our matrices, yielding a new matrix as a result.

The first is **scalar multiplication**. Regardless of the size or shape, if the matrix is not empty we multiply the scalar value by every entry in the matrix, yielding a new matrix. We do this for every entry in the matrix.

The second is **addition**. The shape of the two matrices *must be the same* for addition to go forward. If the shapes are the same and they are both not empty, we add the same row/col element of each argument matrix into the same row/col element of a new matrix, yielding the new matrix. We do this for every element in the two matrices.



Scalar Multiplication



Addition of two matrices

## Requirements

We provide a `lab07_functions.h` As always, you will submit a `lab07_functions.cpp` to Mimir for testing. Write your own `main` to test your code.

We will use a `vector<vector<long>>` as the underlying representation of our matrix. This means that the top level vector has, as elements, another vector.

In `lab07_functions.h` we provide two using definitions to make things a little easier, to wit:

```
using matrix_row = vector<long>;
using matrix = vector<matrix_row>;
```

This is really a big win! We need only say that the type of vector in a `matrix_row` is a `long` and then, if we are careful, can easily change the type of our entire code set by just changing that one template.

## Function Declarations

The functions are clearly described in the `lab07_functions.h` file provided, read them there.

## Printing

I think printing the 2D matrix is actually kind of hard. Here are some tips to help out:

In the include file `io manip` is an io manipulator `setw`. It sets the width for an output element:

- Unlike every other manipulator, it requires you to run it each time you use it.
- If you say something like `cout << setw(5) << 123;` then 5 spaces are reserved for output, 3 of which are occupied by 123 and two of which are just blank spaces (the default, you can change that with `setfill`)
- Two other manipulators are `left` and `right` for left or right justification respectively. Thus
  - `cout << right << setw(5) << 123;` prints 2 spaces and 123
  - `cout << left << setw(5) << 123;` prints 123 and 2 spaces
- If you use an `ostringstream` (and you should), then any `endl` in the stringstream counts as a character in the stream. That might matter when you try to match up with Mimir output.
- My code used:
  - `ostringstream` to capture the output and then convert to a string
  - `setw` to set the width, where the default is 3
  - `right` to get the elements right justified so they look better

## Other Hints

1. Write your own local `lab07_main.cpp` so you can test your code. It would be a good idea to starting doing your testing locally!

- a. As a side note, I see far too many people *writing all the code before testing*. That's crazy! Write a function, test a function, write the next function, test that function, etc. This is the way you figure things out, one by one.
2. You can make a temporary row (of type `matrix_row`) and `push_back` values on to that. You can then `push_back` the row onto a matrix (of type `matrix`). You can reuse the row in the your loop, but remember to `.clear()` it first.
3. Testing is on Mimir. Here is a change! Some of the test cases are hidden so you can't hardcode to the test. By hidden I mean that you can see if you passed but not what the input/output pairs are. The cases are provided as pairs:
  - a. The first of the pair you can see the input output testing
  - b. The second of the pair, you can see if you passed the test but not the input/output pair.